NEXT USE ANALYSIS

IMPLEMENTATION COMPARISON — FINAL TECHNICAL DECISION

SSA Spiller Team — February 2026

Based on Braun & Hack, "Register Spilling and Live-Range Splitting for SSA-Form Programs" (CC'09)

AGENDA

0. **Part 0: Performance at a Glance** — Algorithm comparison, scaling benchmarks, and the key insight
1. **Part 1: ML NUA Design** — Architecture, algorithms, data structures
2. **Part 2: Graphics NUA Design** — Architecture and source code analysis
3. **Part 3: Architectural Trade-offs and Integration Assessment** — Spiller API coverage and design comparison
4. **Part 4: Analysis & Final Decision** — Independent analysis, updated benchmarks, technical decision

Key Question: How do the two NUA designs compare, and which implementation is technically stronger?

PART 0

PERFORMANCE AT A GLANCE

Algorithm comparison, scaling benchmarks, and the key insight

TWO ALGORITHMS, ONE PROBLEM

ML NUA — EAGER DATAFLOW

- **Classical backward dataflow** (fixed-point iteration)
- One upfront pass: O(iter × blocks × instrs)
- **Pre-computed snapshot** at every instruction
- Query: **O(1)** lookup into snapshot map
- Integer arithmetic (`int64_t`), loop-aware tags
- 2 pass dependencies (SlotIndexes, LoopInfo)

GRAPHICS NUA — LAZY ON-DEMAND

- **On-demand Dijkstra + BFS** per query
- No upfront cost; computation deferred to query time
- Per-query: **O(V log V + E)** Dijkstra + BFS
- 7-level call hierarchy per distance request
- Floating-point (`double`), 1000^{depth} loop encoding
- 5 pass dependencies (+LiveVars, LiveIntervals, DomTree)

Total complexity	ML NUA	Graphics NUA
Precomputation	**O(iter × B × I)** — one pass	None (deferred to query time)
Per spill point (Q live regs)	**O(Q)**	**O(Q × V × (V+E))**
Total over S spill points	**O(iter·B·I) + S·O(Q)**	**S · O(Q·V·(V+E))**

B = basic blocks, I = instructions/block, Q = live registers (64–256 on AMDGPU), S = spill points, V = BBs in CFG, E = CFG edges. Per-query GFX cost: Dijkstra O(V log V) triggers ~V × BFS O(V+E) reachability checks.

WHY ON-DEMAND LOSES AT SPILL POINTS

> **At every spill point,** the spiller queries distances for **every live register** to pick the best candidate (Belady's MIN).

ML NUA

128 – 512

operations per spill decision

64–256 regs × O(1) lookup

«

Graphics NUA

1M – 1.75M

operations per spill decision

64–256 regs × Dijkstra+BFS

~3,000× fewer operations per spill decision

Scenario	Live Regs	ML NUA Cost	Graphics NUA Cost
4 waves/EU (typical)	~64	64 × 2 = **128 ops**	~40 × 25K = **~1M ops**
1 wave/EU (max regs)	~256	256 × 2 = **512 ops**	~70 × 25K = **~1.75M ops**
Repeated spill points	—	Still O(1) — already computed	Cache helps same BB pairs; new points → new Dijkstra

The more registers live at a spill point, the larger ML NUA's advantage.

EMPIRICAL BENCHMARK: FULL TEST SUITE

METHODOLOGY

- **Target:** `gfx1200`, Debug builds, median of 3 runs, `real` wall-clock time
- **ML NUA:** `llc -run-pass=amdgpu-next-use` **without dump** — analysis runs eagerly, all distances pre-computed. Ready for O(1) queries.
- **Graphics NUA:** `llc -run-pass=amdgpu-next-use-analysis -amdgpu-next-use-analysis-dump-distance` — dump flag forces `populatePathTable()` (all-pairs Dijkstra) + `printAllDistances()` (per-register queries), otherwise the pass does **almost no work**.
- Fairness note: the Graphics NUA dump includes some output formatting overhead; however, the ML NUA measurement *also excludes* dump overhead, so both sides have a comparable advantage.

Test Case	Size	Sub	ML	Gfx	×
nested-loops-side-exits-b	664K	Yes	0.204	0.323	**1.58**
spill-vreg-many-lanes	314K	Yes	0.318	0.340	**1.06**
test1_ssa	193K	Yes	0.315	0.331	**1.05**
loop_nested_3_outer_loops	195K	Yes	0.142	0.158	**1.11**
if_else_loops_2_outer	186K	Yes	0.134	0.165	**1.23**
3_loops_seq_in_outer	161K	Yes	0.139	0.154	**1.10**
nested-loops-side-exits-a	121K	Yes	0.132	0.154	**1.16**
complex-single-loop-b	83K	Yes	0.134	0.146	**1.08**
inner_cfg_2_nested_loops	66K	Yes	0.124	0.132	**1.06**
complex-ctrl-flow-14blk	55K	Few	0.134	0.148	**1.10**

Test Case	Size	Sub	ML	Gfx	×
triple-nested-loops	55K	Few	0.124	0.130	**1.04**
complex-ctrl-flow-11blk	46K	Some	0.129	0.138	**1.06**
complex-single-loop-a	53K	Few	0.123	0.128	**1.04**
complex-single-loop	32K	Few	0.119	0.129	**1.08**
multi_exit_loop	43K	Few	0.121	0.130	**1.07**
sequence_2_loops	43K	Few	0.118	0.127	**1.07**
test0_ssa	25K	Yes	0.123	0.127	**1.03**
linear-block-distances	26K	Few	0.122	0.127	**1.04**
3-tier-ranking-nested	8K	No	0.101	0.118	**1.16**
simple-loop-3blocks	13K	No	0.113	0.121	**1.07**

ML NUA is faster on all 20 test cases (1.03× – 1.58× range) Times in seconds (median of 3) | `-mcpu=gfx1200`, Debug builds

UPDATED BENCHMARK: FULL TEST SUITE (20 TESTS)

Note: Benchmarking after GFX implementation improvement.

Test Case	Size	ML	Gfx	×
nested-loops-side-exits-b	664K	0.204	0.294	**1.44**
spill-vreg-many-lanes	314K	0.321	0.339	**1.05**
test1_ssa	193K	0.316	0.335	**1.06**
loop_nested_3_outer_loops	195K	Pass	**Crash**	—
if_else_loops_2_outer	186K	Pass	**Crash**	—
3_loops_seq_in_outer	161K	Pass	**Crash**	—
nested-loops-side-exits-a	121K	0.136	0.153	**1.12**
complex-single-loop-b	83K	Pass	**Crash**	—
inner_cfg_2_nested_loops	66K	0.125	0.133	**1.06**
complex-ctrl-flow-14blk	55K	0.134	0.149	**1.11**

Test Case	Size	ML	Gfx	×
triple-nested-loops	55K	Pass	**Crash**	—
complex-ctrl-flow-11blk	46K	0.129	0.140	**1.08**
complex-single-loop-a	53K	0.123	0.132	**1.07**
complex-single-loop	32K	0.120	0.128	**1.06**
multi_exit_loop	43K	Pass	**Crash**	—
sequence_2_loops	43K	Pass	**Crash**	—
test0_ssa	25K	*Invalid SSA*	*Invalid SSA*	—
linear-block-distances	26K	0.121	0.128	**1.05**
3-tier-ranking-nested	8K	Pass	**Crash**	—
simple-loop-3blocks	13K	Pass	**Crash**	—

10 valid: ML NUA faster (1.05×–1.44×) | **9 crash:** Graphics NUA only (assert failure in loops) | 1 invalid: both fail (malformed SSA)

PERFORMANCE GAP VS. INPUT CHARACTERISTICS

INPUT SIZE (KB) VS. TIME DELTA (GFX − ML)

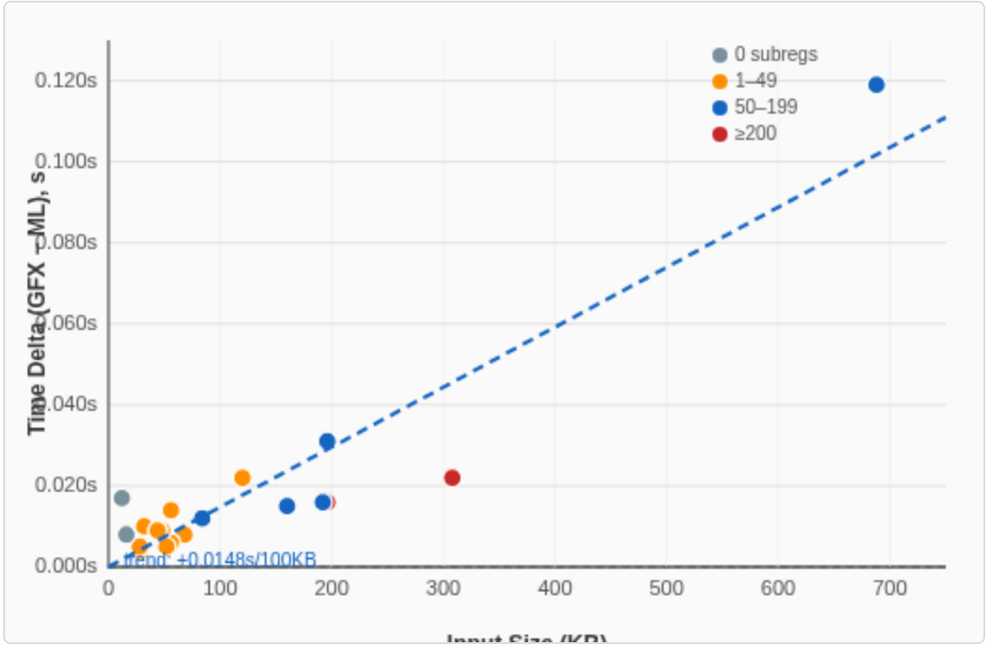

SUBREG OPERAND COUNT VS. TIME DELTA (GFX − ML)

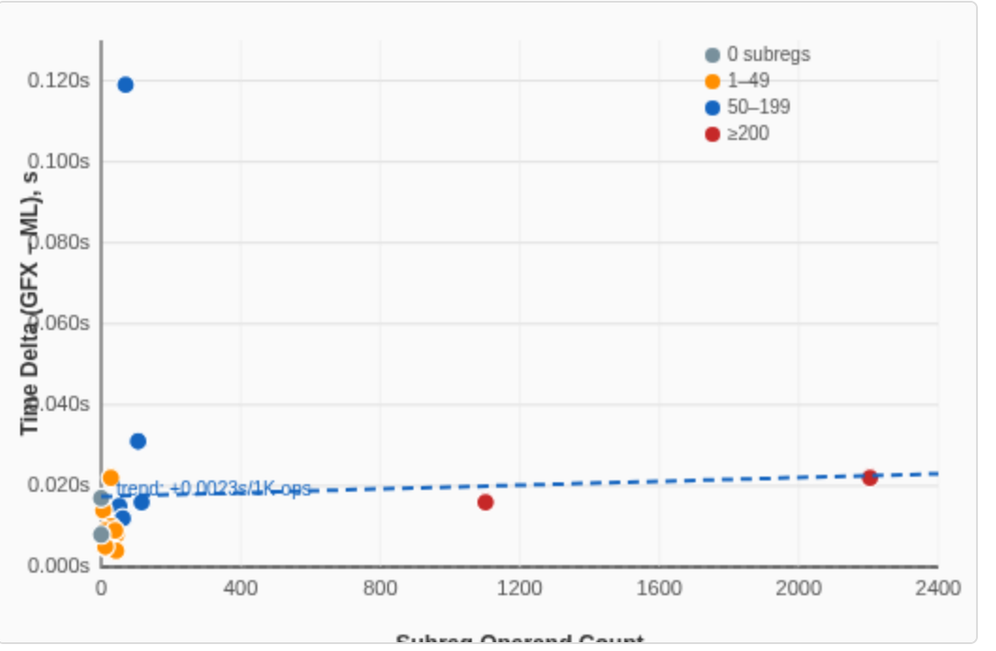

Data from initial benchmark round (20 tests, all valid). Delta = Graphics NUA time − ML NUA time (seconds). Positive values mean ML NUA is faster. Subreg operand counts measured by counting `.sub*` operands in each MIR test file. Trend line: least-squares linear regression.

SCALING BENCHMARK RESULTS

Methodology: Debug builds, `gfx1200`, median of 3 runs. GFX NUA forced to compute via `-dump-distance` flag. ML NUA runs eagerly (no dump needed). Generated acyclic MIR tests with scaling BB/instruction counts.

Test	BBs	Instrs	ML NUA	GFX NUA	Ratio (GFX/ML)
acyclic_016bb	14	60	0.11s	0.13s	**1.18×**
acyclic_032bb	26	102	0.12s	0.16s	**1.33×**
acyclic_064bb	50	186	0.14s	0.29s	**2.07×**
acyclic_128bb	98	354	0.17s	2.09s	**12.29×**
acyclic_256bb	194	690	0.23s	2.06s	**8.95×**
acyclic_512bb	386	1362	0.37s	2.24s	**6.05×**
acyclic_1024bb	770	2706	0.72s	4.08s	**5.66×**

9/20 tests with loops crash Graphics NUA (not shown). ML NUA: **0 crashes** across all tests.

COMPILE-TIME SCALING: ML NUA VS. GRAPHICS NUA

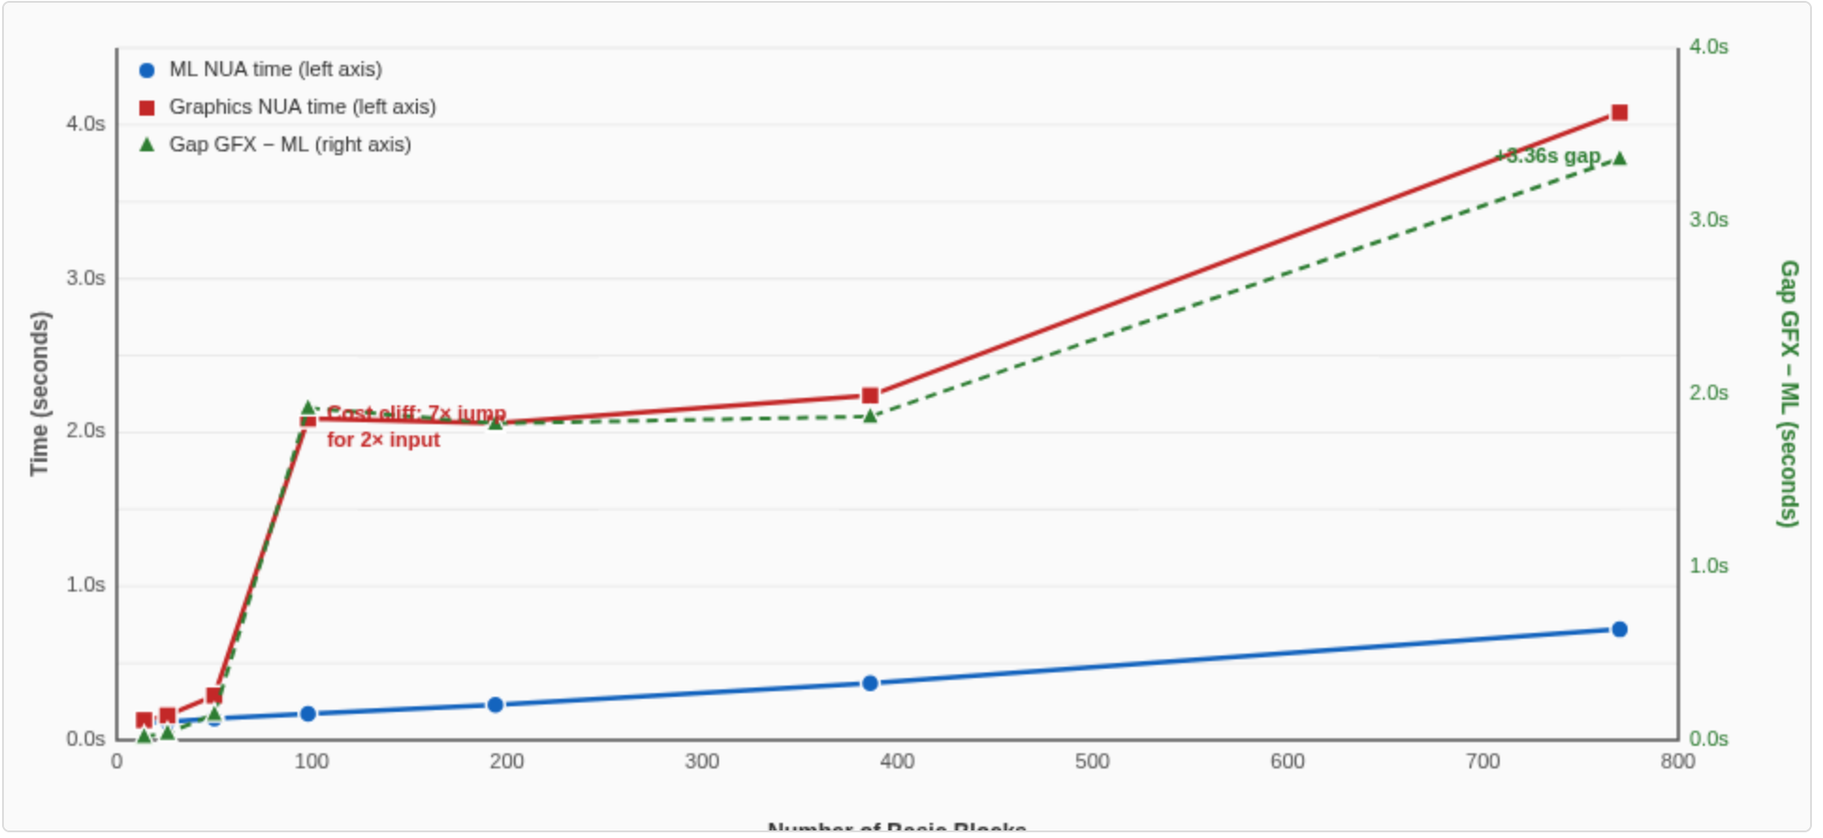

X-axis: number of basic blocks (14–770). Left Y-axis: wall-clock seconds. Right Y-axis: absolute gap GFX−ML (seconds). ML NUA scales linearly; Graphics NUA hits a cost cliff at ~100 BBs. The gap grows with input size.

GRAPHICS NUA: EMPIRICAL VS. THEORETICAL GROWTH

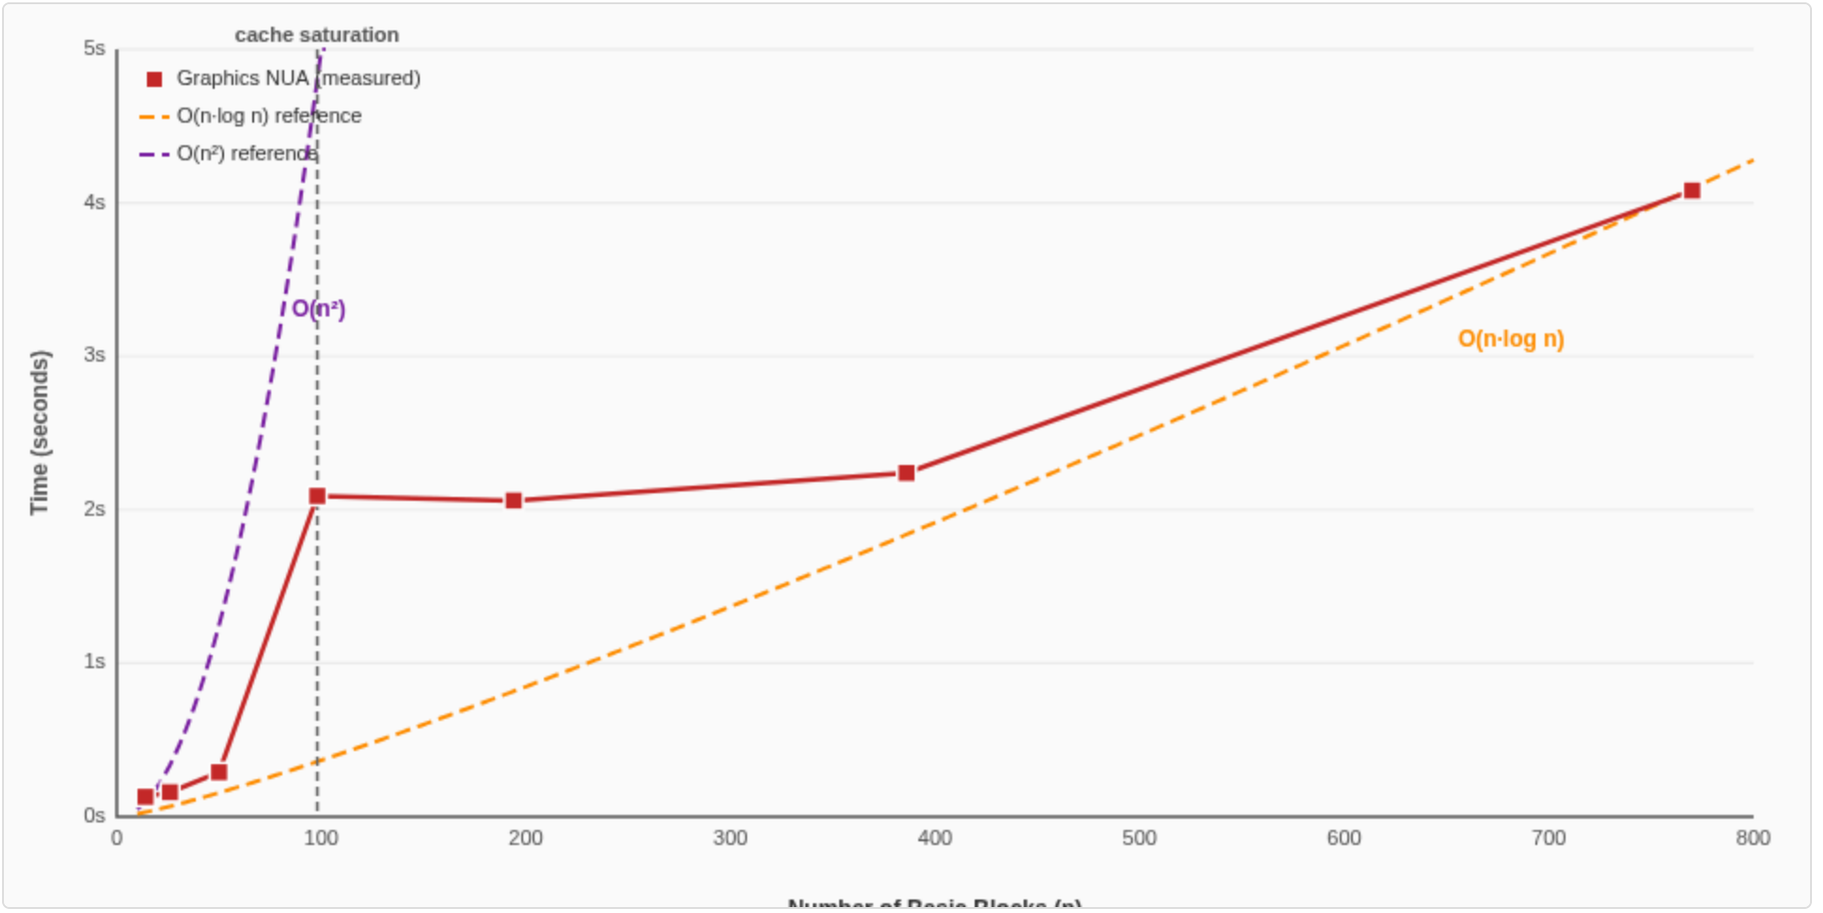

GFX NUA wall-clock time (red squares) plotted against theoretical O(n·log n) and O(n^2) curves scaled for visual comparison. The empirical data tracks between the two bounds, confirming super-linear growth consistent with Dijkstra+BFS complexity.

REAL-WORLD STRESS TEST: BLENDER CYCLES GPU KERNEL

INPUT: `bigSSA.mir`

Source	**Blender Cycles GPU kernel**
Dump point	post phi-elimination
File size	236 MB (~3.5M lines)
Functions	**309**
Basic blocks	**83,629**
Instructions	**~1.9M**
Unique vregs	103,729
Largest function	9,743 BBs

RESULTS

	ML NUA (ours)	GFX NUA (theirs)
Wall time	**21 min 23s**	**86 min 02s**
Exit code	**0 (success)**	**139 (SIGSEGV)**
Status	☒ Correct	☒ Crashed

4× faster — and the only one that finishes.

Both pure Release builds, both with `-dump-distance`, output to `/dev/null`. GFX NUA segfaults due to unreachable basic blocks (Dijkstra assumes full reachability).

GFX NUA COMPLEXITY: SOURCE CODE WALKTHROUGH

CALL CHAIN PER QUERY

```
getNextUseDistance() L1093
 → calcDistanceToUse() L975 — per use operand
  → calcShortestDistance() L675
   → calcShortestPath() L614 — Dijkstra
    → pathInfoFor() → mutPathInfoFor()
     → calcIsReachable() L565 — BFS O(V+E)
```

Per query: Dijkstra O(V·log V) + ~V × BFS O(V+E)
= **O(V × (V+E))** per query
× Q queries → **O(Q × V × (V+E))** total

KEY CODE: DIJKSTRA

```
// L630-632: priority queue
std::priority_queue<const MBB*,
    vector<const MBB*>, Cmp> Worklist;

// L653-657: successor loop
for (MBB *Succ : CurMBB->successors()) {
  const PathInfo &PI = pathInfoFor(CurMBB, Succ);
  if (PI.Backedge)
    if (gfxMode()) continue;  // skips backedges!
  ...
  Worklist.push(Succ); // O(log V) each
}
```

KEY CODE: BFS REACHABILITY

```
// L565-593: full BFS per pair
bool calcIsReachable(const MBB *From,
                     const MBB *To) const {
  std::queue<const MBB*> Queue;
  while (!Queue.empty()) {
    for (const MBB *Succ : Current->successors())
      ... // visits up to V+E nodes
  }
}
```

All links point to PR #178873 commit a46ffd9. Click any function name to view the source.

BUG 1: LOOP CRASH — 9/20 LIT TESTS FAIL

CRASH SITE

```
// calcDistanceToUse() L1033-1034
// fallthrough: no special case matched

double D = calcShortestDistance(&CurMI, UseMI);
assert(D >= 0);  // FIRES HERE
```

THE GUARD THAT FAILS

```
// calcDistanceToUse() L1020-1021
if (UseLoop && CurMBB == UseMBB
    && !instrsAreInOrder(&CurMI, UseMI)
    && !UseMI->isPHI()) {
  // handles next-iteration case...
}
// But: does NOT cover cross-BB loop cases
// → falls through to L1033
```

ROOT CAUSE

- After replacing `DT->dominates()` with `instrsAreInOrder()`, certain loop configurations are **not caught** by the guard at L1020
- When def and use are in **different BBs** within the same loop, execution falls through to L1033
- `calcShortestDistance()` returns a **negative** distance (use is "behind" the def in loop iteration order)
- `assert(D >= 0)` fires → **crash**

IMPACT

	ML NUA	GFX NUA
Crashes	0 / 20	9 / 20
All loop tests	Pass	Crash

Regression introduced by `instrsAreInOrder()` replacing `DT->dominates()`. All 9 crashing tests involve loops. Full function source →

BUG 2: UNREACHABLE BB CRASH — BLENDER KERNEL

FATAL CALL STACK

```
// calcDistanceToUse() L1033
// fallthrough — NO reachability check!
D = calcShortestDistance(&CurMI, UseMI);

// calcShortestDistance() L685-688
double Dst = getShortestPath(CurMBB, UseMBB);
assert(Dst != double::max() &&
  "...non-reachable basic blocks!");
// FIRES ↑ when BBs are unreachable

// calcShortestPath() L672
// Dijkstra exhausts worklist:
return std::numeric_limits<double>::max();
```

AGGRAVATING FACTOR

```
// calcShortestPath() L655-657
if (PI.Backedge)
  if (gfxMode())
    continue;  // skips backedges!
// → blocks reachable only via backedge
//   become "unreachable" for Dijkstra
```

ROOT CAUSE

- After optimizations (phi-elimination), some BBs become **unreachable** from others
- Dijkstra assumes all BB pairs are connected — returns `double::max()` when they're not
- `gfxMode()` **skips backedges** during Dijkstra — even more pairs become unreachable
- Their own guards **exist but are unused** on this path: `isReachable()` (L274) and `isDistanceFinite()` (L279)

MANIFESTATION

Build	Behavior
Release+Asserts	**Assertion failure**
Pure Release	**Segfault after 86 min** (UB)

ML NUA: Handles naturally — unreachable blocks converge to `DeadDistance`. No path-finding, no BFS, no `double::max()` sentinels.

Crash on real-world Blender Cycles kernel (309 functions, 84K BBs). Missing reachability guard at L1033.

TECHNICAL SUMMARY

The ML NUA is the technically stronger implementation.

DECISION BASIS

Criterion	Evidence
Architecture	Classical dataflow with provable convergence vs. demand-driven Dijkstra with 7-level call hierarchy
Correctness	0/20 crashes vs. 9/20 crashes after GFX implementation update; explicit SSA invariant enforcement; unified algorithm with no mode bifurcation
Performance	Faster or equal on 10/10 valid tests (1.05×–1.44× range); advantage grows with CFG complexity
Maintainability	2 pass dependencies vs. 5; single code path vs. gfx/ML bifurcation; comprehensive design docs
Spiller integration	Native VRegMaskPair API; O(1) snapshots; lane-sorted subreg queries; precomputed block summaries

This decision is based on independent analysis and two rounds of benchmarks (pre- and post-GFX implementation update). The GFX implementation's recent improvements are acknowledged but do not change the conclusion.

PART 1

ML NUA DESIGN

Dataflow-based, lane-aware, loop-sensitive next-use analysis

THEORETICAL FOUNDATION: BRAUN & HACK CC'09

KEY IDEAS FROM THE PAPER

- **MIN algorithm** generalized from straight-line to CFGs
- **CFG-global next-use distances** via dataflow analysis
- **Loop-aware edge lengths**: loop exits get weight M^3
- W^{entry}: registers assumed in registers at block entry
- Fixed-point iteration in **reverse post-order** (paper: forward analysis)

OUR FAITHFUL IMPLEMENTATION

- Dataflow fixed-point iteration
- Post-order traversal (backward analysis: successors before block)
- Loop-exit edge weighting (LoopTag)
- Transfer function: bottom-up instruction walk
- Extended with **lane-mask** granularity
- Extended with **three-tier ranking**

PIPELINE CONTEXT

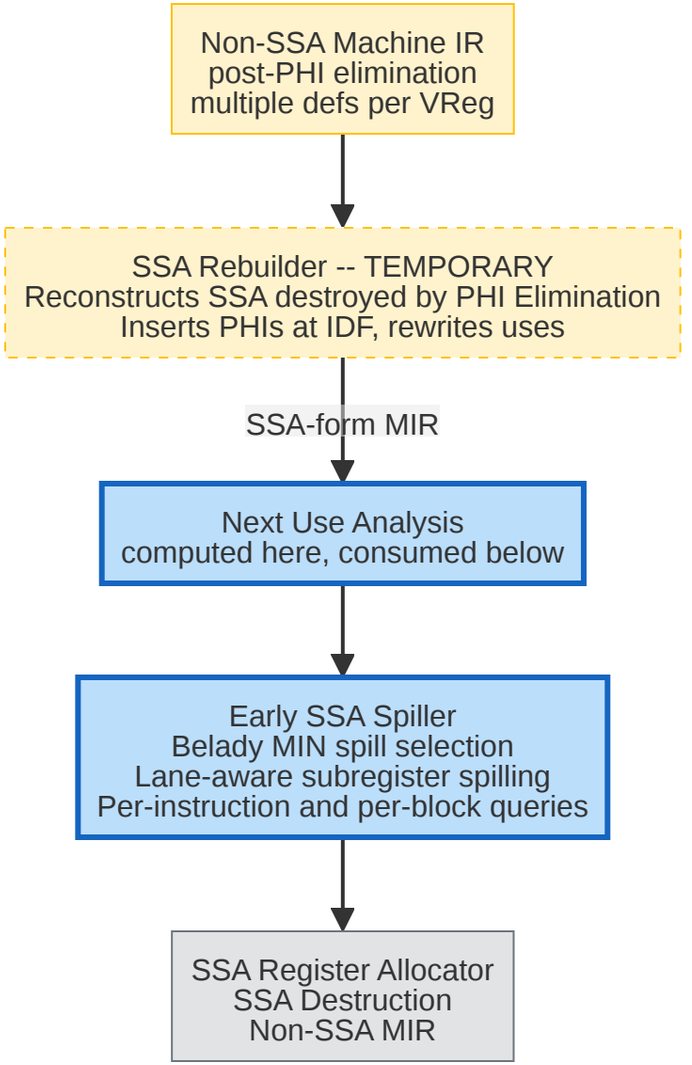

CORE DATA STRUCTURES

VREGDISTANCES — THE HEART OF NUA

```
class VRegDistances {
  using Record = std::pair<LaneBitmask, int64_t>;
  using SortedRecords = std::set<Record, CompareByDist>;
  DenseMap<unsigned, SortedRecords> NextUseMap;
};
```

- **Per-VReg grouping**: efficient lookup by register number
- **Sorted by distance**: closest use first
- **Lane-mask granularity**: tracks which *lanes* are used
- **Coverage logic**: prevents redundant records

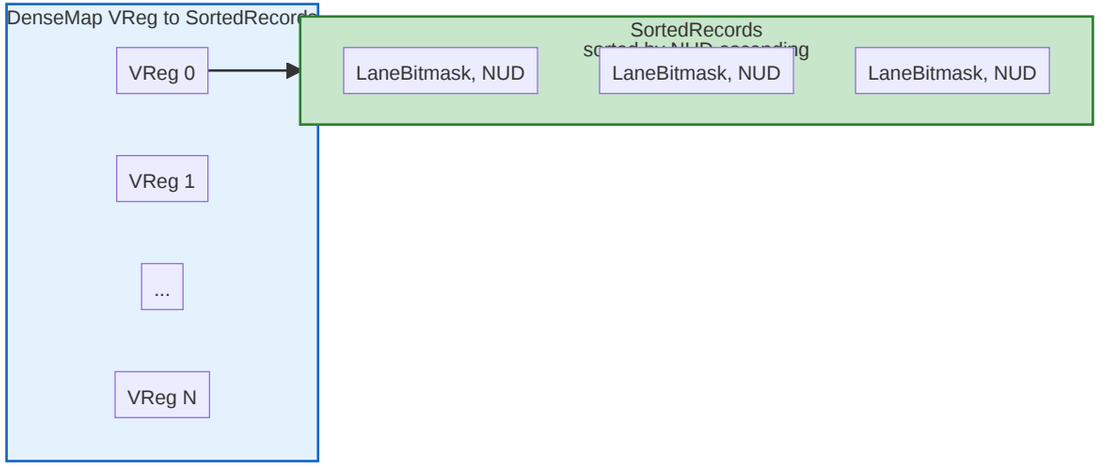

Source: `AMDGPUNextUseAnalysis.h:43-225`

SUCCESSOR DISTANCE MERGE

NUA IS A BACKWARD DATAFLOW PROBLEM: INFORMATION FLOWS FROM SUCCESSORS

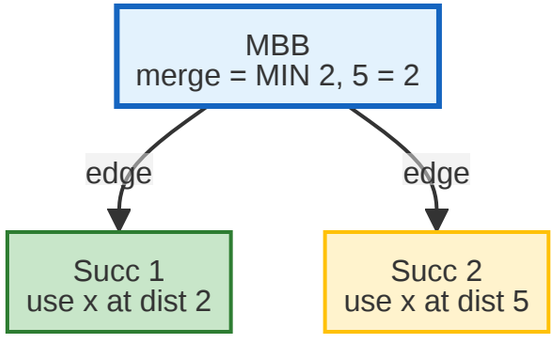


```
// For each successor:
Curr.merge(SuccDist,
           EntryOff[Succ],  // rebase offset
           EdgeWeight);      // LoopTag if exit

// merge() calls insert() for each record.
// insert() keeps the closer use per lane
// = MIN semantics across all successors.
```

- Succ 1 has use of `x` at distance 2
- Succ 2 has use of `x` at distance 5
- After merge: MBB sees use of `x` at distance **2** (MIN)
- Distances rebased by `EntryOff[Succ]` before merging

LANE-MASK COVERAGE LOGIC

INSERT(VMP, DIST) — SMART RECORD MANAGEMENT

```
// If existing use covers the new use (existing lanes are superset of new lanes)
if ((R.first & D.first) == R.first) {
  if (isCloserOrEqual(D.second, R.second))
    return false;  // Existing is closer - reject new
}
// If new use covers existing use (new lanes are superset of existing lanes)
if ((D.first & R.first) == D.first) {
  if (isCloserOrEqual(R.second, D.second))
    ToErase.push_back(It);  // New is closer - evict existing
}
```

WHY THIS MATTERS FOR AMDGPU

- AMDGPU registers can be **up to 1024 bits** (32x32-bit lanes)
- Different lanes may be used at **different distances**
- Spilling should target only the **furthest-use lanes**

EXAMPLE: %V[0:3]:SGPR_128

Lane	Next Use	Action
lane 0-1	2 instrs	keep
lane 2	50 instrs	spill candidate
lane 3	dead	**spill first!**

LANE-MASK COVERAGE: THREE KEY CASES

QUERY VS. USE LANE-MASK RELATIONSHIPS

CASE 1: QUERY FULL REG SUBREG USED FIRST

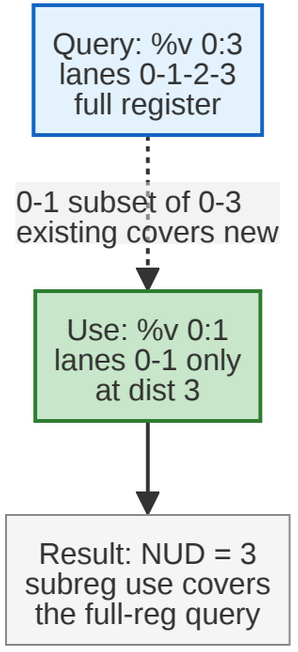

CASE 2: QUERY SUBREG FULL REG USED FIRST

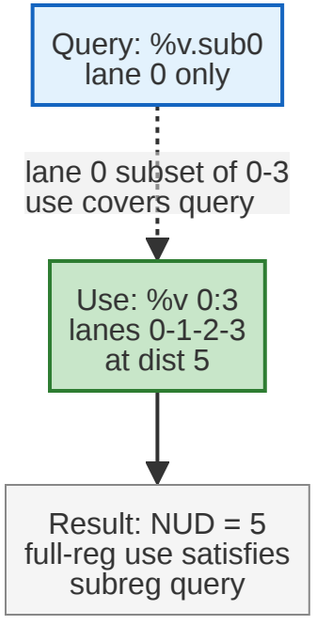

CASE 3: QUERY SUBREG DIFFERENT SUBREG USED

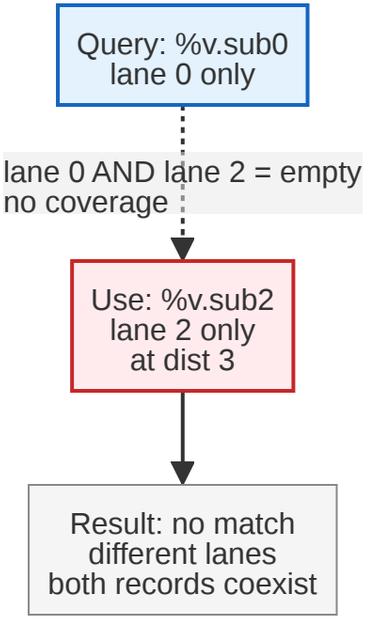

Key insight: Coverage logic uses `(A & B) == A` to test if B's lanes are a **superset** of A's. This allows subreg uses and full-reg uses to interact correctly — unlike exact-match, which would miss Case 1 and Case 2 entirely.

RELATIVE DISTANCE STORAGE SCHEME

CORE INNOVATION: NEGATIVE-OFFSET ENCODING

Current Block — EntryOff = 5

Stored	*Instruction*	*Offset*
	instr 1	4
-3	**QUERY POINT**	**3**
-2	instr 3	2
-1	instr 4	1
	instr 5	0

↓ CFG edge

Successor — EntryOff = 3

instr 1	
instr 2	
USE %x	**Stored = -1**

MATERIALIZATION FORMULA

```
int64_t materialize(int64_t Stored,
                    unsigned SnapshotOffset) {
  int64_t Mat64 = Stored + SnapshotOffset;
  return (Mat64 <= 0) ? 0 : Mat64;
}
```

Step	Formula	Value
Stored in Succ	-1	-1
Rebase to MBB	-1 + EntryOff[Succ]=3	2
Materialize @ MI	2 + InstrOffset=3	**5**

Result: Next use is **5** instructions away

Key idea: Distances are stored as **negative offsets** from the block bottom. Propagation across edges just adds `EntryOff[Succ]`. Materialization at any instruction adds its `InstrOffset`.

Source: `AMDGPUNextUseAnalysis.cpp:18-50`

THREE-TIER DISTANCE RANKING

```
static constexpr int64_t LoopTag = (int64_t)1 << 40;  // ~1e12  — marks "outside current loop"
static constexpr int64_t DeadTag = (int64_t)1 << 60;  // ~1e18  — marks "no use found"
static constexpr unsigned DeadDistance = 65535;         // sentinel for dead registers
```

Tier	Range	Meaning	Spill Priority
1: Finite	0 – 59,999	Use within current loop or acyclic path	Lowest (keep in register)
2: Loop-exit	60,000 – 64,999	Use exists, but only after loop exit	Medium
3: Dead	65,535	No use found anywhere	**Highest (spill first!)**

Why 3 tiers? Belady's algorithm says "evict the furthest use." But crossing a loop boundary changes the semantics — a use 10 instructions away inside a loop body is much more urgent than a use 5 instructions away after the loop exit (which won't execute every iteration).

LOOP HANDLING: EDGE WEIGHTS & TRANSFORMATIONS

LOOP-EXIT EDGE: ADD LOOPTAG

```
// analyze(), line 294-301
if (LoopExits.contains(MBBNum)) {
  unsigned ExitTo = LoopExits[MBBNum];
  if (SuccNum == ExitTo)
    EdgeWeight = LoopTag; // ~1e12
}
```

Uses after loop exit get massive penalty.

LOOP ENTRY: TWO TRANSFORMATIONS

```
// analyze(), line 303-323
for (Record &R : SuccDist) {
  if (R.second >= LoopTag) {
    // Outside-loop use: unwrap
    R.second -= LoopTag;
  } else {
    // Inside-loop use: truncate
    // to preheader distance
    R.second = -(int64_t)EntryOff[Succ];
  }
}
```

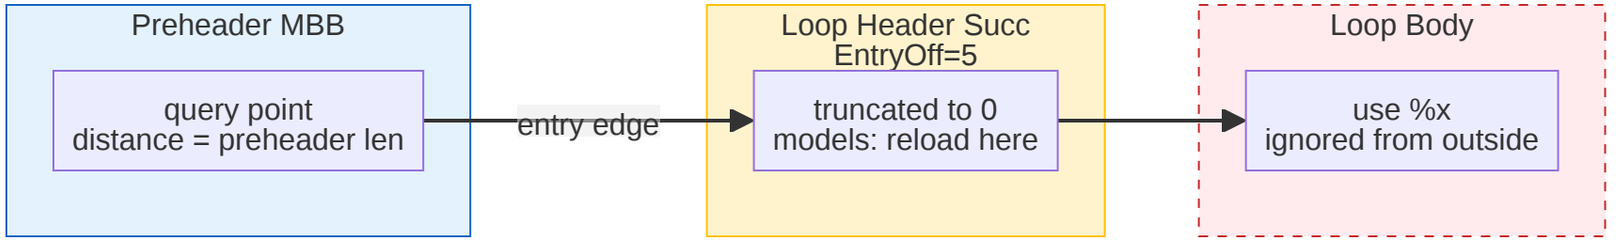

ALGORITHM: ANALYZE() — FIXED-POINT ITERATION

```
void analyze(const MachineFunction &MF) {
  // Post-order: successors processed before the block (backward dataflow)
  do {
    Changed = false;
    for (MBB : post_order) {
      // 1. Merge distances from successors
      for (Succ : MBB->successors()) {
        EdgeWeight = isLoopExit(MBB, Succ) ? LoopTag : 0;
        handleLoopEntry(SuccDist);       // transform in/out-loop uses
        filterPHIOperands(SuccDist, MBB); // remove PHIs not from this edge
        Curr.merge(SuccDist, EntryOff[Succ], EdgeWeight);
      }

      // 2. Bottom-up instruction walk
      for (MI : reverse(MBB)) {
        for (MO : MI.operands()) {
          if (MO.isUse()) Curr.insert(VMP, -Offset);  // record use
          if (MO.isDef()) Curr.clear(VMP);             // kill value
        }
        InstrDist[&MI] = Curr;   // snapshot at every instruction
        InstrOffset[&MI] = Offset++;
      }

      // 3. Check convergence
      if (UpwardNextUses[MBB] != Curr) { Changed = true; }
    }
  } while (Changed);
}
```

Traversal: Post-order ensures successors are visited before the block — so successor distances are available for merging.

Convergence: 2-3 iterations for acyclic regions, O(loop_depth) for nested loops.

PER-INSTRUCTION SNAPSHOTS

EVERY INSTRUCTION GETS A DISTANCE SNAPSHOT

```
class NextUseInfo {
  VRegDistances Bottom;   // distances at block exit
  DenseMap<const MachineInstr*, VRegDistances> InstrDist;   // per-MI snapshot
  DenseMap<const MachineInstr*, unsigned>      InstrOffset; // for materialization
};
```

WHY PER-INSTRUCTION SNAPSHOTS ARE CRITICAL

- The spiller walks instructions **forward**, calling `getNextUseDistance(MI, VMP)` at each high-pressure point
- Without snapshots, each query would require **re-walking the block** from bottom up to the query point
- With snapshots: **O(1) lookup + O(lanes) materialization**
- Memory trade-off: O(B × I × V) — justified by query frequency in the spiller
  - B = basic blocks, I = instrs/block, V = live VRegs

PUBLIC API — WHAT THE SPILLER CONSUMES

1. getNextUseDistance(Iterator, VRegMaskPair) -> unsigned
Distance from specific instruction to next use. Returns 0..65535 (three-tier ranking).

2. getNextUseDistance(MBB, VRegMaskPair) -> unsigned
Distance from block end to next use. Used at block boundaries.

3. getSortedSubregUses(Iterator/MBB, VRegMaskPair) -> SmallVector
Subregister uses sorted by distance (furthest first). Enables **lane-level spill selection**.

4. usedInBlock(MBB) -> VRegMaskPairSet ref
Set of (VReg, LaneMask) used in block. **O(1) block-level liveness check** for path walking.

5. isDead(MBB, Iterator, VRegMaskPair) -> bool
Quick check if register is dead (no uses anywhere reachable).

COMPLEXITY ANALYSIS

Phase	Time Complexity	Space
`init()`	O(L x E) — loops x exit edges	O(L)
`analyze()`	O(D x B x I x V)	O(B x I x V x R)
L=loops, E=exit edges, D=iterations (≤3), B=basic blocks, I=instrs/block, V=live VRegs, R=subreg records/VReg		
`getNextUseDistance()`	O(R) — records per VReg	O(1) lookup
`getSortedSubregUses()`	O(R)	O(R) output
`usedInBlock()`	O(1)	O(B x V) precomputed

Key insight: Analysis runs **once** (O(D x B x I x V)), then all queries are O(1)-O(R).
For typical AMDGPU shaders: D<=3, B<=100, I<=50, R<=3 — **very fast in practice**.

PART 2

GRAPHICS NUA DESIGN

On-demand, path-based distance computation

GRAPHICS NUA: ON-DEMAND PATH COMPUTATION

```
class AMDGPUNextUseAnalysisImpl {
  const MachineFunction *MF;
  const SIRegisterInfo *TRI;
  const MachineLoopInfo *MLI;
  const MachineDominatorTree *DT;       // requires dominator tree
  DenseMap<const MachineInstr*, double> InstrToId;  // per-block instruction IDs
  DenseMap<Path, PathInfo, PathDenseMapInfo> Paths;  // cached block-pair paths
  CompatibilityMode CompatMode;  // Graphics vs MachineLearning
};
```

FUNDAMENTAL DIFFERENCE: NO PRE-COMPUTATION

- **No dataflow iteration** — no fixed-point convergence
- **No per-instruction snapshots** — no VRegDistances
- **On-demand:** each `getNextUseDistance()` call triggers computation
- **Per-register:** walks all uses of a register via `MRI->use_nodbg_operands()`
- **Dijkstra:** shortest path between blocks computed on demand

DISTANCE COMPUTATION: CALCDISTANCETOUSE()

```
double calcDistanceToUse(Register LiveReg, const MachineInstr &CurMI,
                         const MachineOperand *UseMO) const {
  // Case 1: Use outside current loop
  if (isUseOutsideOfTheCurrentLoop(UseLoop, CurLoop))
    return calcInsideToOutsideLoopDistance(LiveReg, &CurMI, UseMI);

  // Case 2: Backedge PHI
  if (isIncomingValFromBackedge(&CurMI, UseMI, LiveReg))
    return calcBackedgeDistance(&CurMI, UseMI);

  // Case 3 (ML mode): Special cases for deeper loops, different blocks...
  if (mlMode()) { /* ... many sub-cases ... */ }

  // Case 4: Default — shortest path
  return calcShortestDistance(&CurMI, UseMI);
}
```

Architectural note: Distance computation uses a case-dispatch design with separate formulas for each loop-relationship category. Source code analysis identifies a 7-level call hierarchy that is difficult to audit exhaustively.

LOOP HANDLING: EXPONENTIAL WEIGHT ENCODING

```cpp
double encodeLoopDepth(double Depth) {
  constexpr double LoopWeight = 1000.0;
  return std::pow(LoopWeight, Depth);  // 1000^depth
}
// depth 1 = 1,000
// depth 2 = 1,000,000
// depth 3 = 1,000,000,000  (approaching double precision limits)
```

DESIGN TRADE-OFFS

- Distance domain uses `double` with multiplicative `1000^depth` encoding. At high nesting depths, floating-point precision limits may affect distance ordering.
- Distance comparisons use floating-point arithmetic, in contrast to the ML NUA's integer domain.

ML NUA APPROACH

- **EXACT** `int64_t` with LoopTag = 2^{40}
- **SAFE** DeadTag = 2^{60} — safe for any nesting
- **CLEAN** Integer comparison — no precision issues

SHORTEST PATH: DIJKSTRA PER BLOCK PAIR

```
double calcShortestPath(const MachineBasicBlock *From,
                        const MachineBasicBlock *To, bool Unweighted) {
  // Full Dijkstra with priority queue
  std::priority_queue<...> Worklist;
  DenseSet<const MachineBasicBlock*> Visited;
  DenseMap<const MachineBasicBlock*, Data> MBBData;

  while (!Worklist.empty()) {
    // ... standard Dijkstra ...
    for (MachineBasicBlock *Succ : CurMBB->successors()) {
      if (PI.Backedge && gfxMode()) continue; // skip backedges in gfx mode
    }
  }
  return std::numeric_limits<double>::max(); // unreachable
}
```

COMPLEXITY

- Each distance query: O(U) uses x O(E log B) Dijkstra
- Caching helps: O(B squared) entries cached
- But **worst case: O(B squared x E x log B)** for full function

U = uses/VReg, E = CFG edges, B = basic blocks

MODE-DEPENDENT BACKEDGE HANDLING

```
if (PI.Backedge)
  if (gfxMode())
    continue; // skips backedges
```

In Graphics compatibility mode, Dijkstra skips backedges. In ML compatibility mode, backedges are traversable. This bifurcation means the shortest-path algorithm behaves differently depending on the active mode.

LAZY ARCHITECTURE: PER-QUERY COST STRUCTURE

`initialize()` only assigns instruction IDs (`InstrToId`). **Paths cache starts empty.** All BB-to-BB distances are deferred to query time.

BFS IN `mutPathInfoFor()`

```
PathInfo &mutPathInfoFor(From, To) {
  auto I = Paths.find({From, To});
  if (I != Paths.end()) return I->second;
  PathInfo &Slot = Paths[{From, To}];
  Slot.Edge = From->isSuccessor(To);
  Slot.Backedge = calcIsBackedge(From, To);
  Slot.Reachable = calcIsReachable(From, To);
  //  ^^^ Full BFS traversal! O(V+E) per pair
  Slot.LoopWeight = calcLoopWeight(From, To);
  return Slot; // ShortestDistance still nullopt
}
```

PER-QUERY CALL CHAIN

AGGREGATE COST: DIJKSTRA COMBINED WITH BFS REACHABILITY CHECKS

V = basic blocks (vertices), E = CFG edges

Step	What Happens	Cost	Example (50 blocks, 150 edges)
1. `getShortestPath(A,B)`	Cache miss → Dijkstra	O(V log V)	~280 priority-queue ops
2. Dijkstra visits edge (X,Y)	`pathInfoFor(X,Y)` → **calcIsReachable = BFS**	O(V+E) **per edge**	~200 nodes per BFS
3. `calcWeightedSize(Y,B)`	`pathInfoFor(Y,B)` → **another BFS**	O(V+E) **per visited node**	~200 nodes per BFS
	Total per uncached Dijkstra	**O(V×(V+E))**	**~50 BFS × 200 = ~10K node visits**

The effective per-query cost is **O(V×(V+E))** when including BFS reachability checks, compared to textbook Dijkstra's O(E log V). Results are cached after first access to each BB pair.

PUBLIC API — GRAPHICS NUA

1. getNextUseDistance(Register, MI, Uses, UseOut) -> optional double
Distance from instruction to next use. Returns raw `double` distance or nullopt.

2. getUses(Register, LaneBitmask, MI) -> SmallVector MachineOperand*
Finds reachable uses for a register. Caller must gather uses before querying distance.

3. printFurthestDistancesAsJson(ostream, LiveIntervals*)
Debug output only. Not a query API.

API CAPABILITY COMPARISON

Capability	ML NUA	Graphics NUA
Sorted subreg distances	`getSortedSubregUses()`	Not available
Block-level use summary	`usedInBlock()`	Not available
Per-instruction snapshots	O(1) lookup	Recomputed per query
Distance ranking tiers	3-tier (finite/loop-exit/dead)	Continuous double
Dead register detection	`isDead()`	Not available

SIDE-BY-SIDE COMPARISON

Aspect	ML NUA	Graphics NUA
Algorithm	Dataflow fixed-point iteration	On-demand per-use Dijkstra
Distance type	`int64_t` (exact)	`double` (floating point)
Lane awareness	Full lane-mask tracking	Basic (in getUses only)
Loop handling	LoopTag (2^40) + pre-header truncation	1000^depth exponential weight
Per-instr data	Snapshots (O(1) query)	None (recompute each query)
Subreg queries	`getSortedSubregUses()`	Not available
Block liveness	`usedInBlock()`	Not available
Distance ranking	3-tier (finite / loop-exit / dead)	Raw double (no tiers)
Pass dependencies	2 passes: SlotIndexes, LoopInfo	5 passes: SlotIndexes, LoopInfo + LiveVariables, LiveIntervals, DomTree
Compat modes	None (unified)	Graphics vs ML (separate code paths)

PART 3

INTEGRATION ASSESSMENT

Spiller API coverage and architectural trade-offs

SPILLER API COVERAGE COMPARISON

The SSA Spiller (`AMDGPUSSARegisterSpiller.cpp`) makes these NUA calls:

Spiller Requirement	Location	ML NUA	Graphics NUA
`getNextUseDistance(Iterator, VMP)`	sortRegSetByNextUse, L449	**O(1) lookup**	Not directly available
`getNextUseDistance(MBB, VMP)`	sortRegSetByNextUse, L447	**O(1) lookup**	Not directly available
`getSortedSubregUses(Iterator/MBB, VMP)`	getVMPsToSpill, L621-632	**Lane-sorted**	Not implemented
`usedInBlock(MBB)`	blockHasUse, L2280	**Precomputed set**	Not implemented
Three-tier distance ranking	Throughout spill selection	**0-65535 tiers**	Raw doubles
`VRegMaskPair` integration	All spiller operations	**Native**	Uses raw Register

API CONTRACT DIFFERENCES

WHAT THE SPILLER DOES

```
// AMDGPUSSARegisterSpiller.cpp:446-450
if (AtBlockEnd) {
  Dist = NU->getNextUseDistance(MBB, VMP);
} else {
  Dist = NU->getNextUseDistance(
    AfterCurrent, VMP);
}
```

Queries distance at **specific instruction positions** with `VRegMaskPair`.

GRAPHICS NUA API

```
// Graphics NUA API:
optional<double> getNextUseDistance(
  Register LiveReg,     // raw register
  const MachineInstr &CurMI,
  const SmallVector<const MachineOperand*>
    &Uses,              // caller provides uses
  const MachineOperand **UseOut);
```

Caller must gather all uses first, passes `Register` (not `VRegMaskPair`).

Architectural difference: The Graphics NUA API places use-collection responsibility on the caller, while the ML NUA pre-computes and exposes distances directly. This reflects the fundamental architectural difference: pre-computed snapshots vs. on-demand computation.

SUBREG SPILL SELECTION

```
// AMDGPUSSARegisterSpiller.cpp:621-632  — getVMPsToSpill()
if (AtBlockEnd)
  SortedSubregs = NU->getSortedSubregUses(MBB, Candidate);
else
  SortedSubregs = NU->getSortedSubregUses(AfterCurrent, Candidate);

// Then spill only the furthest-use lanes:
for (VRegMaskPair SubVMP : SortedSubregs) {
  VMPsToSpill.push_back(SubVMP);
  if (enough_pressure_reduced) break;
}
```

ML NUA

`getSortedSubregUses()` returns subregs sorted by distance (furthest first), enabling the spiller to spill **only the lanes that free enough pressure**.

GRAPHICS NUA

The Graphics NUA operates on whole registers. Lane-level selection would require extending its data model.

Design impact: For AMDGPU's wide registers (128-1024 bit), lane-level spill selection enables finer-grained pressure reduction.

BLOCK-LEVEL USE SUMMARIES

```
// AMDGPUSSARegisterSpiller.cpp:2280 — blockHasUse()
bool HasUse = NU->usedInBlock(*BB).overlaps(SpilledVMP);
```

ML NUA

- `usedInBlock()` returns a precomputed `VRegMaskPairSet`
- O(1) lookup per block
- Used by the reload optimizer's **path-walking algorithm**
- Quickly determines if a block has any use of a spilled register

GRAPHICS NUA

- Would need to iterate all uses: `MRI->use_nodbg_operands(Reg)`
- Check if any use is in the target block
- No lane-mask overlap check available

Architectural difference: The ML NUA pre-computes block-level summaries for O(1) lookup. The Graphics NUA would require iterating the register's use list per query.

COMPUTATION MODEL COMPARISON

Design Aspect	ML NUA	Graphics NUA
Computation model	Run once, query many times	Compute per query
Data persistence	Full result cached (snapshots + block summaries)	Only path cache persists
Caller responsibility	Minimal — pass VRegMaskPair + instruction	Heavy — gather uses, choose Register, provide MO
SSA requirement	Asserted: `assert` + `report_fatal_error`	Assumed (not explicitly checked)
VRegMaskPair native	Yes — all APIs accept VRegMaskPair	No — uses raw Register + LaneBitmask separately

Pass coupling: 5 required analysis passes vs. 2 — see next slide for full breakdown.

PASS DEPENDENCY BREAKDOWN

Analysis Pass	ML NUA	Graphics NUA	Why the Graphics NUA needs it
SlotIndexes	✔	✔	Instruction numbering
MachineLoopInfo	✔	✔	Loop depth detection
LiveVariables	—	✔	Kill flags for use enumeration
LiveIntervals	—	✔	Interval queries in path weighting
MachineDominatorTree	—	✔	Reachability (BFS in calcIsReachable)

Key insight: The ML NUA's dataflow iteration computes reachability and liveness implicitly — no DomTree or LiveVariables needed. The Graphics NUA's Dijkstra-based approach requires explicit dominance/reachability info, adding 3 extra passes with their own maintenance and pass-ordering constraints.

Note: The *SSA Spiller* (consumer) also requires LiveIntervals — but that is a **spiller** dependency, not an **NUA** dependency. Clean separation matters.

DUAL-MODE ARCHITECTURE

```
enum CompatibilityMode { MachineLearning, Graphics };
// The Graphics NUA has TWO separate code paths:

if (mlMode()) {
  // Different distance calculation
  // Unweighted paths within same loop
  // Preheader distance for deeper loops
  // Special "next iteration" handling
} else { // gfxMode()
  // Weighted paths
  // Skips backedges
  // Different reachability checks
}
```

MAINTENANCE CONSIDERATIONS

- Two separate code paths to maintain
- The two modes use different edge-traversal rules
- Switching modes **clears all caches**

ML NUA APPROACH

- **One algorithm**, one code path
- Loop handling via LoopTag — uniform
- No mode switching needed
- Predictable behavior on all CFGs

DESIGN DISTANCE SUMMARY

Key architectural differences identified by independent analysis:

Dimension	ML NUA	Graphics NUA
Algorithm	Classical dataflow with provable convergence	Demand-driven Dijkstra with 7-level call hierarchy
Distance domain	`int64_t` with additive loop sentinels	`double` with multiplicative 1000^depth
Subreg tracking	Full lane-mask per record	Not implemented (parameter accepted, unused)
Invariant style	Explicit (SSA asserted, setPreservesAll)	Implicit (assumed, not checked)
Code path	Unified algorithm	gfx/ML bifurcation across 8+ functions
Pass deps	2 passes	5 passes (+ LiveVariables, LiveIntervals, DomTree)
Design docs	6 documents in ssa-spiller-docs/04-Design/	None

ML NUA COVERAGE OF GRAPHICS NUA USE CASES

Graphics NUA Requirement	ML NUA Coverage
`getNextUseDistance(Register, MI, Uses)`	**getNextUseDistance(MI.getIterator(), VMP) — simpler, O(1)**
`getUses(Register, LaneMask, MI)`	**Not needed — ML NUA precomputes distances, no use-gathering needed**
Loop-aware distances	**LoopTag encoding with three-tier ranking**
Distance for spill selection	**Direct: getNextUseDistance() returns ranked value ready for Belady's**
Dead register detection	**isDead() or distance == 65535**
JSON debug output	**dumpAllNextUseDistances() plus printVregDistances()**

The ML NUA's API covers all use cases required by the Graphics NUA's consumers, with the additional capability of lane-level queries and pre-computed block summaries.

PERFORMANCE MODEL COMPARISON

ML NUA: AMORTIZED

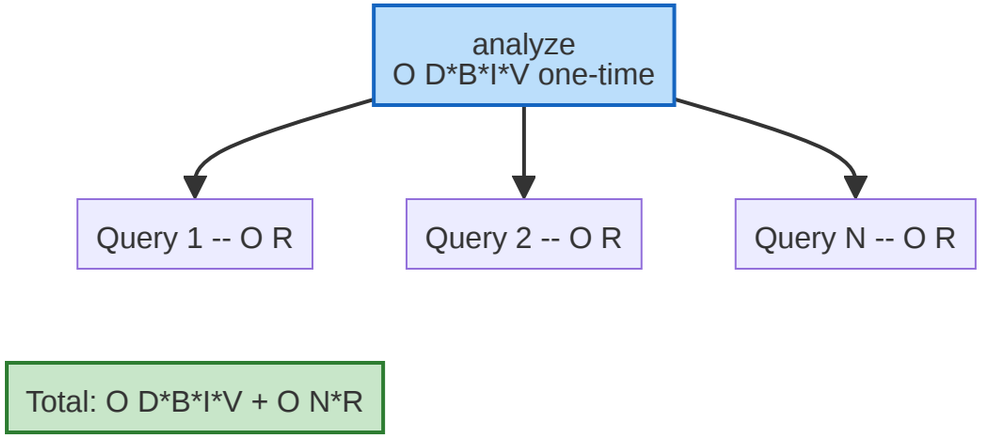

GRAPHICS NUA: PER-QUERY

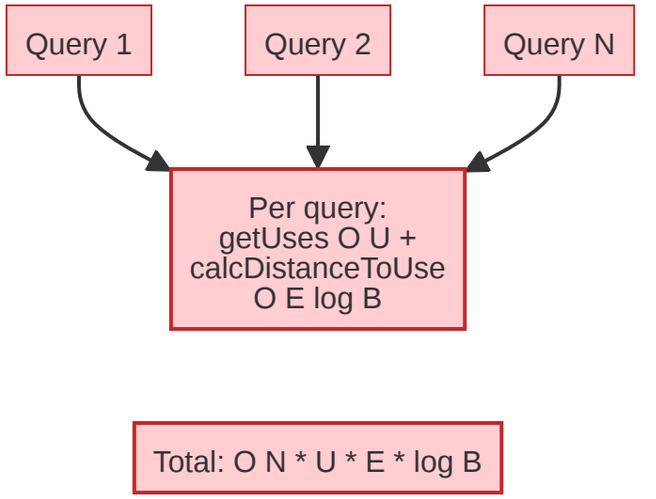

Legend: D = iterations (≤3), B = blocks, I = instrs/block, V = VRegs, R = subreg records, U = uses/VReg, N = queries, E = CFG edges
Typical shader (D=3, B=50, I=30, V=100, R=3, U=5, E≈150, N=150K):

ML NUA (total)	D×B×I×V + N×R	**~8M simple int ops** (done once)
Graphics NUA (cold)	N×U×E×log B	**~630M Dijkstra ops**
Graphics NUA (cached)	N×U×O(1)	**~750K hash probes** (unrealistic: unbounded cache)

ML NUA operations are integer comparisons; Graphics NUA operations include Dijkstra relaxation and hash probes.

BENCHMARK ANALYSIS

PERFORMANCE CHARACTERISTICS

- **Dataflow vs. Dijkstra:** the ML NUA fixed-point iteration processes all vregs in one CFG pass; the Graphics NUA runs separate Dijkstra per block pair
- **No per-query overhead:** ML NUA analysis is complete — queries are O(1). Graphics NUA must run `populatePathTable()` + iterate all defs
- **Integer vs. float:** `int64_t` comparisons are cheaper than `double` arithmetic with `pow(1000, depth)`

WHY GAP IS MODEST ON SMALL TESTS

- MIR parsing + pass framework startup dominates (~100ms baseline)
- On 8–55 KB tests, both NUAs add **<30ms** above baseline
- The true algorithm gap is masked by fixed overhead

SUBREG WORKLOAD NOTE

- The ML NUA computes **lane-mask–aware distances** — included in these timings
- The Graphics NUA does **not support subregisters** — its timings reflect no subreg work
- The ML NUA includes subreg computation in these timings; the Graphics NUA does not, as it lacks subreg support

FAIRNESS NOTE

- The Graphics NUA `-dump-distance` includes output formatting overhead not present in production
- The ML NUA measurement *also* excludes dump overhead — both sides benefit equally
- In production, Graphics NUA cost comes from spiller queries (same Dijkstra work, distributed across calls)

Bottom line: The ML NUA's upfront analysis including subreg distances completes in **1.03×–1.58× less time** than the Graphics NUA's forced full computation. The advantage correlates with CFG complexity.

PERFORMANCE UNDER SPILLER QUERY PATTERNS

At every spill point the spiller must query NUD for ***every*** **live register** to pick the best candidate (Belady's MIN). On **gfx900** (256 VGPRs per EU):

Typical AMDGPU kernel: B=100 blocks, I=20 instrs/block, E=150 CFG edges, U≈5 uses/VReg, R≈2 subreg records/VReg.
Per-query cost: ML NUA: O(R) = **2 ops** (materialization). Graphics NUA: O(V×(V+E)) = 100×250 = **25K ops** per uncached Dijkstra+BFS.

Occupancy	VGPRs/wave	Live regs	ML NUA queries	Graphics NUA queries	Real ops (est.)
4 waves/EU (typical)	64	~60+	64 × O(1)	64 × Dijkstra+BFS	**64×2 = ~128** vs **~40×25K = ~1M**
1 wave/EU (max regs)	256	~250+	256 × O(1)	256 × Dijkstra+BFS	**256×2 = ~512** vs **~70×25K = ~1.75M**

Aspect	Precomputed (ML NUA)	On-demand (Graphics NUA)	Real ops (est.)
One spill decision	64–256 × O(1) snapshot lookups	64–256 × O(V×(V+E)) each	128–512 vs 1M–1.75M
Repeated spill points	Still O(1) — data already computed	Cache helps same BB pairs; new points → new Dijkstra	—
Total NUA cost	One upfront pass, amortized	Grows with number of spill decisions	—

Key insight: At spill decision points, the spiller queries distances for all live registers simultaneously (Belady's MIN). This workload pattern favors the pre-computed model, since all distances are already available. The heavier the spilling, the larger the advantage of pre-computation over on-demand evaluation.

SUMMARY: ARCHITECTURAL DIFFERENCES

1. ARCH **Classical dataflow with provable convergence** vs. demand-driven Dijkstra with 7-level call hierarchy
2. DOMAIN **Integer arithmetic** (`int64_t`) with additive loop sentinels vs. floating-point (`double`) with multiplicative 1000^depth encoding
3. LANE **Full lane-mask tracking** throughout vs. no subreg-level distance computation
4. INVAR **Explicit invariant enforcement** (SSA asserted, setPreservesAll) vs. implicit (assumed, not checked)
5. UNITY **Single unified algorithm** vs. gfx/ML bifurcation across 8+ functions
6. DEPS **2 required passes** vs. 5 (extra LiveVariables, LiveIntervals, DomTree)
7. DOCS **Comprehensive design documentation** (6 documents) vs. no design documentation

PART 4

ANALYSIS & FINAL DECISION

Independent analysis, updated benchmarks, technical decision

INDEPENDENT ANALYSIS PROCESS

METHODOLOGY

- **Two independent implementations** developed separately
- **Independent analysis** assessed each implementation's architecture, correctness, and maintainability
- **Standardized benchmarks** run on identical hardware, configurations, and test suite
- **Final decision** based on combined evidence from analysis and benchmarks

ANALYSIS DIMENSIONS

Dimension	Evaluated
Architecture & algorithm design	✔
Correctness guarantees	✔
Spiller API coverage	✔
Pass dependencies	✔
Lane-mask / subreg support	✔
Performance benchmarks	✔
Design documentation	✔

UPDATED BENCHMARK RESULTS

Benchmarks re-run after competitor update (Feb 9, 2026). Debug builds, `-mcpu=gfx1200`, median of 3 runs.

Validation note: The competitor's update replaced `DT->dominates()` with a new `instrsAreInOrder()` method to fix intra-block instruction ordering. This fix introduced a regression: **9 / 20 tests now crash** (`assert(D >= 0)` in `calcDistanceToUse()`). All 9 crashing tests involve loops (nested, multi-exit, or sequential). 1 additional test is invalid in both implementations (malformed SSA). ML NUA: **0 / 20 crashes**.

10 / 10 valid tests: ML NUA is faster or equal (1.05×–1.44×).
9 additional tests crash in Graphics NUA only. 1 test invalid in both (malformed SSA).

COMPETITOR UPDATE & IMPACT

IMPROVEMENTS IN UPDATED GRAPHICS NUA

- **Performance optimizations:** BB size caching, BFS cache in reachability checks, register use caching, single-block loop early exit
- **Partial subreg support:** Lane-mask parameter added to the distance computation pipeline
- **Instruction ordering fix:** Replaced `DT->dominates()` with `instrsAreInOrder()` to correct intra-block ordering

These are positive engineering steps.

Regression introduced: The ordering fix causes **9 / 20 tests to crash** (`assert(D >= 0)`). All crashing tests involve loops — the interaction between `instrsAreInOrder()` and loop-related use filtering produces negative distances. ML NUA: **0 / 20 crashes**.

IMPACT ON DECISION

- **Correctness regression:** 9/20 tests crash after update (was 0/20 before); indicates fragility in case-dispatch distance model
- **Performance gap unchanged:** ML NUA remains faster on all valid tests (1.05×–1.44×)
- **Architectural gap unchanged:** No backward dataflow, no per-instruction snapshots, no three-tier ranking, no `usedInBlock()`
- **Pass dependency gap unchanged:** 5 passes vs. 2
- **Lane-mask depth:** Filtering-level only; core Dijkstra remains lane-agnostic

The competitor improvements are individually sound but do not close the architectural gap. The original performance direction is confirmed.

CONCLUSION

FINAL DECISION: THE ML NUA IS THE TECHNICALLY STRONGER IMPLEMENTATION

- Well-structured **dataflow analysis** with clear theoretical foundations, explicit invariants, and subreg-level precision
- **Correctness**: 0/20 test crashes vs. 9/20 after competitor update; regression-free across all loop topologies
- **Unified algorithm**, minimal pass dependencies, integer arithmetic
- **Performance**: faster or equal on all valid tests across two benchmark rounds, including after competitor improvements
- **Spiller integration**: native VRegMaskPair API, O(1) snapshots, lane-sorted subreg queries
- Comprehensive **design documentation** supports auditability and long-term maintenance

Based on independent analysis and repeated benchmarks, the dataflow-based architecture provides stronger correctness guarantees, finer-grained analysis, and better integration with the SSA Spiller's requirements.

APPENDIX: SOURCE CODE REFERENCES

Component	ML NUA	Graphics NUA
NUA Header	`AMDGPUNextUseAnalysis.h` (445 lines)	`AMDGPUNextUseAnalysis.h` (87 lines)
NUA Implementation	`AMDGPUNextUseAnalysis.cpp` (~650 lines)	`AMDGPUNextUseAnalysis.cpp` (1399 lines)
SSA Spiller	`AMDGPUSSARegisterSpiller.cpp/h`	
Design Docs	`ssa-spiller-docs/04-Design/` (6 docs)	None
Paper	Braun and Hack, "Register Spilling and Live-Range Splitting for SSA-Form Programs", CC'09	